

COMP9517: Computer Vision

2026 T1 Lab 1 Specification

Maximum Marks Achievable: 2.5

This lab is **worth 2.5% of the total course mark**.

Deadline for submission is Week 3, Friday 6 March 2026, 18:00 (Sydney time).

Objective: This lab revisits important concepts covered in the Week 1 and Week 2 lectures and aims to make you familiar with implementing specific algorithms.

Software: You are required to use OpenCV 3+ with Python 3+ and submit your code as a Jupyter notebook (see coding and submission requirements below). In the tutor consultation session this week, your tutors will give a demo of the software to be used, and you can ask any questions you may have about this lab.

Materials: The images to be used in this lab are available on WebCMS3 (under Lab 1).

Submission: Submit your source code as a Jupyter notebook (.ipynb file) that includes all output and answers to all questions (see Coding Requirements below) by the above deadline. The submission link is available on WebCMS3 (under Lab 1).

Coding Requirements

Make sure that in your Jupyter notebook, the input images are readable from a location specified as an argument, and all output images and other requested results are displayed in the notebook environment. *All cells in your notebook should have been executed so that the tutor/marker does not have to execute the notebook again to see the results.*

Task 1 (0.75 mark)

While walking in UNSW campus, you decide to take a nice picture, only to discover that your camera is defective and produces very noisy images. You remember from the computer vision course that you can still create a nice picture simply by **Averaging**. So, you take 10 images from the same viewpoint (see the given .zip file on WebCMS3).


In your notebook, calculate the average image of the 10 noisy images and display the result. Also answer the following questions in your notebook:

Theoretically, how much noise reduction should you be able to achieve this way? That is, by

what factor should the standard deviation of the noise drop?

Practically, how much noise reduction did you actually achieve here? Note that this requires measuring the standard deviation of the pixel values within a region containing only random noise and no other image structures. The best region for this is the sky.

Measure the standard deviation within some rectangular region in the sky in one of the single noisy images and within the same region in the averaged image. Report the two standard deviations and the ratio of the former to the latter.



Task 2 (0.75 mark)

Given the average image, you can also experiment with image filters. One interesting filter is the **Difference of Gaussians (DoG)**.

Let us define two approximations of 2D Gaussian filters:

$$h_1 = \frac{1}{16} \cdot \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 2 & 1 & 0 \\ 0 & 2 & 4 & 2 & 0 \\ 0 & 1 & 2 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix} \quad h_2 = \frac{1}{256} \cdot \begin{bmatrix} 1 & 4 & 6 & 4 & 1 \\ 4 & 16 & 24 & 16 & 4 \\ 6 & 24 & 36 & 24 & 6 \\ 4 & 16 & 24 & 16 & 4 \\ 1 & 4 & 6 & 4 & 1 \end{bmatrix}$$

This notation means that, to get the actual kernel values, each value in the matrix must be

multiplied by the factor in front of it. This is because the kernels must be normalized (the kernel values must sum to 1 to avoid blowing up the pixel values in the filtered image).

As you can see, h_1 is a 'narrower' filter (the values in the outer rows and columns are all 0, so they do not actually count, and the filter is effectively 3 x 3 pixels) and h_2 a 'wider' filter (5 x 5 pixels). These two filters can be combined to create a DoG filter:

$$\text{DoG} = h_1 - h_2$$

The convolution of an image f with the DoG filter can be written as:

$$g = f * \text{DoG} = f * (h_1 - h_2) \quad (1)$$

In other words, you first calculate a new DoG kernel, being the difference ($h_1 - h_2$), and then you convolve the image f with it to get the result image g .

From the image processing lecture, you remember that convolution is a distributive operation. This means the convolution could equivalently be written as:

$$g = f * h_1 - f * h_2 \quad (2)$$

In other words, you first convolve the image f with h_1 , and separately with h_2 , and then you subtract the two resulting filtered images to get g .

Take your average image from Task 1 and calculate its DoG-filtered version using the above two different approaches (1) and (2). *Keep in mind that both approaches involve subtraction, and thus the resulting pixel values could be negative, so make sure to use data types that can handle this correctly.* In your notebook, display the result images of the two different DoG-filtering approaches.

Also answer the following question in your notebook:

While the two approaches achieve the same objective, they offer different advantages. Could you discuss the advantage of each approach relative to the other one?

Task 3 (1 mark)

From the lecture on image processing in the spatial domain, you know that high frequencies can be used to sharpen an image using a technique called **Unsharp Masking**.

Take one of the DoG-filtered images from Task 2 (it does not matter which one) and multiply it with a factor of 5. Then add the resulting image to your average image from Task 1.

You may need to adjust the brightness/contrast of the final image a bit to make the intensities look similar to those of the average image from Task 1.

Visually comparing the output image of Task 3 with the output image of Task 1, you will notice that indeed edges look a bit sharper now (higher local contrast across the image). In your notebook, show the two images side by side, and answer the following questions:

In addition to edge sharpness, what else has been amplified in the output image of Task 3, which we were trying to get rid of in Task 1?

And what is the intuitive explanation for why this happens? Hint: consider the effect of DoG filters on image structures of different frequencies.

Copyright: UNSW CSE COMP9517 Team. Reproducing, publishing, posting, distributing, or translating this lab assignment is an infringement of copyright and will be referred to UNSW Student Conduct and Integrity for action.

Released: 23 February 2026